Voronoi Treemaps in D3

Peter Henry University of Washington
phenry@gmail.com

Paul Vines University of Washington
paul.l.vines@gmail.com

ABSTRACT
Blah blah blah

1. INTRODUCTION

Treemaps are a category of visualizations used for displaying hierarchical data. While node-and-edge diagrams are often used for visualizing hierachical structures, treemaps offer some significant advantages. Primarily, treemaps are space-filling, and therefore allow each node in a hierarchy to have more viewing area devoted to it than in a node-and-edge diagram. This allows both larger hierarchies to be visualized, as well as more detail to be shown on each node, such as additional text, colors, or glyphs to show attributes of the node.

The majority of treemap layouts used are variants of rectangular treemaps. These Have the advantage of being relatively fast to layout and in cases of limited scale produce reasonably understandable treemaps. However, there are three drawbacks to rectangular treemaps.

First, as hierarchies become deeper, the treemap cells can become increasingly extreme in aspect ratio, resulting in narrow rectangles more difficult to see than if their area was distributed in a more square-like space. This problem is mostly mitigated by various tweaks to the treemapping algorithm to try to keep the aspect ratio of regions close to one.

Second, the borders between different regions in the hiearchy can become difficult to see. In particular, two cells neighboring one another in the treemap but not siblings in the hierarchy can appear to share a common edge delineating the same inner node is their parent, when this is in fact not the case. Finally, rectangular treemap algorithms naturally only fill rectangular regions which could be undesirable for aesthetic or practical reasons.

Voronoi treemaps eliminate these problems. Firstly, Voronoi treemap cells are arbitrary polygons but as will be discussed later, the generation algorithm results in generally low aspect ratio cells. Secondly, the fact that Voronoi treemap cells are arbitrary polygons means edges between cells will fall at any angle, rather than only vertical or horizontal, and so two neighboring cells will generally never have a continuous-looking edge unless they are in fact siblings in the hierarchy and thus share the edge of their parent node's cell. Finally, Voronoi treemaps can be produced for any arbitrary polygonal region, and so do not restrict the shape to be filled by the treemap.

Multiple Voronoi treemap algorithms have been created in recent years (CITE). However, none are available for use in a web framework. Our work has been to implement one of the fastest algorithms (CITE) for use in the D3 web framework. Despite the optimizations employed by the algorithm creators, generation of a Voronoi treemap is still a computationally intensive task. Therefore, we have additionally written the D3 module with features to try to allow Voronoi treemaps to be used for web visualizations without causing a poor user experience even on complex datasets.

The remainder of the paper is structured as follows: Section 2 has a discussion of related work including a brief introduction to Weighted Voronoi diagrams and a discussion of the algorithms created for Voronoi treemaps. Section 3 describes the implementation of our work in D3 and optimizations added for client-side web usability. Section 4 shows the use of our framework on several datasets and an evaluation of the computational burden of our system. Section 5 discusses the potential applications of our system. Section 6 concludes with proposals of future work to be done in this space.

2. RELATED WORK
2.1 Voronoi Diagrams

Voronoi diagrams are a technique for dividing a region containing sites into cells to satisfy the following condition: given a distance function $d(p, q)$ where p and q are points, any point p is labeled as belonging to the site q that results in the lowest distance, $d(p, q)$. In this case to be labeled means to be inside a bounding polygon formed for each site. In the case of a simple euclidean distance function, $d(p, q) = \sqrt{(dx^2 + dy^2)}$ this results in a cell border being equidistant between the two closest sites.

For Voronoi treemaps two extensions are made to the basic

Voronoi diagram. First, sites are given initially random positions, a diagram is generated, and then sites are moved to the centroidal positions in their cell and then the diagram is re-generated. This is repeated until a relatively stable set of positions is found (Lloyd). The effect of this iterative process is to create lower aspect ratio cells. Second, rather than using a standard euclidean distance function the generation algorithm uses a weighted distance function, where each site is assigned a weight that corresponds to generating a larger or smaller cell. This allows the sizes of cells to be adjusted to reflect the relative size or number of children of a specific node in the hierarchy being displayed.

After these extensions are made, the Voronoi treemap algorithm proceeds to compute the Voronoi diagram for each level of the hierarchy: it starts at the highest level, generates the Voronoi diagram of the first level of nodes, and then recursively descends into each cell and generates the Voronoi diagram for the children of that node using the cell as the new bounding region. The computational burden of this can be high; several different algorithms for computing the Voronoi diagram have been developed and are briefly summarized below.

2.2 Previous Approaches

Voronoi treemaps have been implemented previously (Balzer) using both additively weighted and geometrically weighted Voronoi diagram algorithms. This implementation used the same iterative algorithm for creating centroidal Voronoi diagrams described above. To create the weighted diagrams, however, it used a sampling algorithm wherein points were sampled in the space and distances to nearest sites calculated, to give an approximation of the correct weighted Voronoi diagram. This results in an algorithm on the order of $O(n^2)$ where n is the number of sites. The benefit of this algorithm is that it the sampling process is the bottleneck and is easily parallelized to achieve linear speedups with additional CPU cores.

This algorithm implementation was improved by Sud et al. (CITE SUD) by using GPU programming to significantly speedup computation by parallelizing across graphics hardware. However, the algorithm remained $O(n^2)$ for the number of sites. Further, this approach is not feasible for web programming because consumer devices are not commonly equipped with powerful graphics cards and do not all support the use of the graphics card by a website (CITE).

The algorithm proposed by Nocaj & Brandes (CITE NOCAJ) offers a significant asymptotic improvement on these previous designs. Rather than a sampling-based approach, this implementation uses the algorithm for computing arbitrary-dimension Power Diagrams proposed by Aurenhammer (CITE AURENHAMMER). In this approach the 2D points representing sites are lifted into 3-dimensional dual space based on their weights. The convex hull made by these 3D points is then computed, and projected back down to 2D to produce the Voronoi diagram. This method is on the order of $O(n \log n)$ and so can provide a significant speedup for generating treemaps of larger datasets.

3. METHODS

The core computational components of our implementation were adapted from a Java implementation of the Nocaj & Brandes algorithm (CITE NOCAJ) using a lift into 3-dimensions followed by computation of the convex hull and projection back into 2-dimensions to create the Voronoi diagram. As with other implementations, we use Lloyd's algorithm to iteratively adjust the site locations to be the centroids of their cells and then adjust the weights of the sites to fit the area of each cell to within an error threshold.

4. RESULTS
5. DISCUSSION
6. FUTURE WORK
7. ACKNOWLEDGMENTS
8. REFERENCES